Framework	Task source	Grading	Leak / hack defense	What we reuse
SWE-bench	Real GitHub PRs + issues	Fail-to-pass unit tests	Grader outside the agent; Verified subset	Real-source tasks, frozen environments
Inspect AI	Dataset + solver + scorer	Rules or model-graded	Sandboxed execution	Task structure, sandbox config
lm-eval-harness	YAML static tasks	Log-likelihood / exact match	n-gram decontamination	Static-QA format, decontamination
BenchJack	— (audit tool)	—	V1–V8 hackability checklist	Pre-release audit